

# Fast Network Threat Detection based on the TRAbID dataset: An Analysis of Various Machine Learning Models

Shrey Agarwal, Ashwini Hiremath, Ibrahim Khalid, Harsh Shinde, Justin Wang

**Abstract**—In our increasingly digital world, network security is paramount. This project aims to enhance threat detection by training various machine learning models to classify network activities as either attacks or normal. Using the TRAbID dataset we trained models including Support Vector Machines, Logistic Regression, Naive Bayes, Random Forest, and XGBoost. These models were evaluated based on various metrics including F1 score and detection time. Results indicated that Gaussian Naive Bayes performed best in terms of F1 score, while Logistic Regression was the fastest. However, many models exhibited overfitting to the known dataset, not performing well with new attack data. This project looks at the challenges in deploying robust threat detection systems and highlights the potential of models like Naive Bayes for such tasks. Limitations included dataset size constraints and computational resources.

**Index Terms**—Network anomaly detection, Machine Learning, Classification, Detection time

## I. INTRODUCTION

IN our ever increasing digital world, it is vital that we also have improved security over our networks. Based on the Cloudflare DDoS 2023 Q4 report [1] and Verizon’s data breach report 2023 [2], there has been an increase in the number of cyber attacks, including network based attacks. By using live network data and classifying it as harmful or not, we can create a system that will improve the peace of mind of a user and allow them to use modern technology with confidence.

The primary objective of this project is to train various machine learning models for the classification of whether an activity on the network is a threat or not. The secondary objective of this project is to evaluate the machine learning models based on their performance, and computation time and cost in order to find the best model. This can be used to create a firewall application that uses the built machine learning model.

## II. LITERATURE REVIEW

Bagui et al. explored the UNSW-NB15 dataset using a hybrid approach that used k-means clustering followed by decision trees or Naive Bayes. Using decision trees, they were able to achieve a 99% accuracy rate. [3] Bhati et al. used the KDDCUP99 dataset to explore the effectiveness of ensemble models such as AdaBoost and Random Forest. [4]

Zhou and Yang developed an “immune algorithm” for detecting network anomalies using SVM. They were able to find that using such an algorithm, they were able to get an

accuracy of 94% with a detection time of 27 seconds for the entire test suite. [5]

Tavallae et al. developed multiple machine learning models trained on the KDDCUP99 dataset, including Random Forest, Naive Bayes, and SVM. They were able to get accuracies of 92.8%, 81.7%, and 65.1%, respectively. [6]

Viegas et al. generated the TRAbID dataset with the goal of providing a better dataset for training machine learning models. They used a novel approach of using various datasets with differing levels of similarity. They trained on this dataset using the Naive Bayes and Decision Tree algorithms. [7]

Based on our short literature review among the listed papers and otherwise, there has not been a clear strive for creating a model with the goal of fast detection after training. We plan on furthering research in this regard by exploring the effectiveness of various machine learning techniques on the TRAbID dataset.

## III. PROPOSED METHODOLOGY

We are planning to use the TRAbID dataset [7] as it provides a relevant dataset for the problem, this can be subject to change. The TRAbID dataset is sufficiently large, with 51 feature columns, including classification, made up of generated network traffic data. Our methodology will be to first explore this dataset using tools such as Pandas, Matplotlib, and SciKit Learn. This includes tasks such as identifying correlations, normalizing values, and selecting important features. Following this, we will work on training a few different types of models in order to determine which models perform the best at distinguishing network threats. Models we plan on training are: Support Vector Machines, Logistic Regression, Naive Bayes, Random Forest, and XGBoost. These selected models can be subject to change. We will compare the models based on accuracy, precision and recall, and speed of classification. The goal of this project is to find a model that, when deployed in a real life scenario, best performs analysis of raw network data and classifies it as ‘normal’ or ‘attack’ in terms of detection time and detection reliability.

## IV. DATA UNDERSTANDING AND EXPLORATION

The dataset was downloaded from the original source as ‘arff’ files. These files were quite large consisting of nearly 28,000,000 rows. Moreover, there was a huge imbalance in the dataset with nearly 99.7% of the entries being ‘normal’ entries. The exact numbers can be seen in Table I.

TABLE I  
CLASS DISTRIBUTION BEFORE AND AFTER REDUCTION

Class	Before Reduction		After Reduction	
	Normal	Attack	Normal	Attack
Known	28,618,341	36,628	48,432	36,628
Similar	28,477,883	10,440	15,969	10,440
New	28,391,902	17,752	58,138	17,752

In order to reduce the computational complexity of this project as well as fix the class imbalance, we used a shell script to reduce the dataset. The results of this class re-balancing and data reduction can be seen in Table I. The reduced files were saved as ‘csv’ files locally for ease of use in data exploration and modelling.

The first thing we need to understand about the data is how the known, similar, and new attacks are related to each other and how they should be used to train and create the model. Based off the paper by Viegas et al. we can see their approach was based off Fig 1.[7] To use the dataset as designed, we modelled our evaluation methods similar to theirs.




Fig. 1. Dataset relations and model evaluations

Now that we know what dataset to use for what purpose, we began exploring the known dataset. There are 51 features in the dataset, all of which are listed in Table II

TABLE II  
LIST OF FEATURES

'ip_type', 'ip_len', 'ip_id', 'ip_offset', 'ip_RF', 'ip_DF', 'ip_MF', 'ip_proto', 'ip_checksum', 'udp_sport', 'udp_dport', 'udp_len', 'udp_chk', 'icmp_type', 'icmp_code', 'icmp_chk', 'tcp_sport', 'tcp_dport', 'tcp_seq', 'tcp_ack', 'tcp_ffyn', 'tcp_fsyn', 'tcp_first', 'tcp_fpush', 'tcp_fack', 'tcp_furg', 'fr_length', 'conn_status', 'count_fr_dst_src', 'count_serv_src_dst', 'num_bytes_src_dst', 'num_bytes_dst_src', 'num_bytes_serv_dst_src', 'num_pushed_dst_src', 'num_syn_fin_src_dst', 'num_fin_src_dst', 'num_fin_dst_src', 'num_ack_dst_src', 'num_syn_src_dst', 'num_rst_dst', 'num_rst_dst_src', 'first_packet', 'first_serv_packet', 'class'
--

One of the first things we noticed was that the columns 'ip\_RF', 'ip\_DF' and 'ip\_offset' were completely empty.

These columns were removed as they should not affect our results greatly. Next, we performed a correlation analysis to see which of these remaining features were strongly correlated with the target feature. The results of this analysis can be seen in Table III. The features with a greater than 0.1 absolute correlation are highlighted in bold.

TABLE III  
SORTED CORRELATION WITH TARGET FEATURE

Feature	correlation	Feature	correlation
class	1.000000	num_syn_dst_src	-0.017306
<b>tcp_sport</b>	0.376711	count_fr_src_dst	-0.018505
<b>first_serv_packet</b>	0.288401	first_packet	-0.032560
<b>tcp_fsyn</b>	0.183862	num_ack_dst_src	-0.042007
<b>tcp_ffyn</b>	0.179369	num_ack_src_dst	-0.055285
<b>udp_sport</b>	0.159150	count_serv_dst_src	-0.058866
<b>num_rst_dst_src</b>	0.148479	icmp_type	-0.076590
tcp_furg	0.081623	tcp_seq	-0.082221
ip_id	0.070370	tcp_first	-0.089140
ip_proto	0.068846	ip_type	-0.098455
tcp_fpush	0.068064	<b>icmp_chk</b>	-0.106911
udp_chk	0.050838	<b>icmp_code</b>	-0.115398
udp_dport	0.043169	<b>num_bytes_serv_src_dst</b>	-0.205258
num_syn_fin_dst_src	0.032562	<b>conn_status</b>	-0.217683
num_syn_src_dst	0.023503	<b>udp_len</b>	-0.244732
num_fin_src_dst	0.018179	<b>fr_length</b>	-0.249022
count_fr_dst_src	0.016557	<b>ip_len</b>	-0.265645
num_syn_fin_src_dst	0.010139	<b>num_bytes_src_dst</b>	-0.275082
num_pushed_src_dst	-0.005405	<b>tcp_dport</b>	-0.302774
num_pushed_dst_src	-0.006834	<b>ip_DF</b>	-0.305169
num_fin_dst_src	-0.008403	<b>tcp_fack</b>	-0.406649
num_rst_src_dst	-0.008925	<b>tcp_ack</b>	-0.456281
count_serv_src_dst	-0.010181	<b>num_bytes_serv_dst_src</b>	-0.520069
ip_checksum	-0.015762	<b>num_bytes_dst_src</b>	-0.535816

For further understanding, we performed principal component analysis (PCA) on all of the features as well as only those features with an absolute correlation greater than 0.1. PCA analysis is one of the common ways to reduce the dimensionality of a given dataset. In the case of machine learning, it means building a smaller model. The benefits of a smaller model are that they are more efficient and faster to train than models using the full dataset. One good benchmark for deciding the number of features to use from PCA reduction is to select the number of features that can help explain 95% of the variance in the dataset. We found that when analyzing all the features present, we needed only 28 of the PCA transformed features to achieve 95% cumulative explained variance. When analyzing only those features that had a greater than 0.1 absolute correlation, we needed only 10 of those PCA transformed features to achieve 95% cumulative explained variance. See Fig. 2

Moreover, since that dataset itself covers various different ranges in different features, it was important that we also perform feature scaling. This was done through the use of the standard scaler class from the scikit learn library.

From our data understanding, it is clear that we can approach this problem via a number of different angles; we can use all the features, or a subset of highly correlated features, or we can use PCA reduced feature spaces, or some combination.

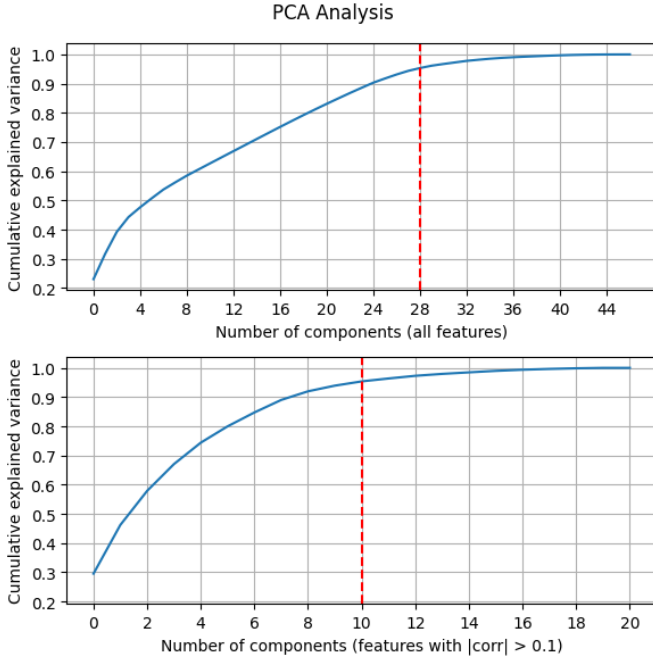


Fig. 2. PCA analysis on all features and features with correlation greater than 0.1. 95% cumulative explained variance line highlighted

The rest of this paper will look at how we used these different angles, or ‘pipelines’, along with our selected models to find the answer to our problem. The pipeline that produces a dataset of all features scaled using the standard scaler will be called ‘*scaled*’. The pipeline that produces a dataset of all those features which had an absolute correlation of 0.1 or more with the target feature and is scaled using the standard scaler will be called ‘*corr, scaled*’. The pipeline that produces a dataset that explains 95% of the variance in all features will be called ‘*PCA*’. The pipeline that produces a dataset that explains 95% of the variance of all those features which had an absolute correlation of 0.1 or more with the target feature will be called ‘*corr, PCA*’.

Since our proposed methodology is trying to find the best model in terms of detection time and detection reliability, it is important to design our modelling to reflect that. While it is straight forward to test the models in terms of various metrics like accuracy and F1 score, it is harder to test their effectiveness in terms of time to detect. We need to develop our models and test suite in such a way that it allows us to do this with ease.

## V. MODELLING

For modelling, we have selected 5 models to test against. These are: Naive Bayes, Support Vector Machines (SVM), Logistic Regression, Random Forest, and XGBoost. These models will be tested in all the pipelines mentioned in Section IV. Moreover, since we are testing the detection time of the models to correctly classify raw data, we need to design a benchmarking function that can cover all such scenarios. The pseudo code for this benchmarking function is as shown as follows:

```

1 def benchmark(df, pipeline_fn, model):
2     df_benchmark = pipeline_fn(df)
3     X = df_benchmark[features]
4     y = df_benchmark[target]
5     y_pred = model.predict(X)
6     save accuracy, precision, recall, F1
7     start = time()
8     for _ in ITER:
9         df_benchmark = pipeline_fn(df)
10        X = df_benchmark[features]
11        model.predict(X)
12    end = time()
13    save (end - start / ITER)

```

Based on this code and the problem at hand, we developed the project architecture as described in Fig. 3. There we can see that all datasets are passed through the same selected pipeline function as defined in Section IV. The ‘Known’ dataset is then split into a 70:15:15 ratio for training, validation, and testing, respectively. Using the training set of the known dataset, we train a baseline model. Then based on that, we perform hyper parameter tuning using the validation set. Once we have the best parameters for the given model, we train a new model on the training set using these parameters. This final model is not only tested on the testing set of the known data, but also on the ‘Similar’ dataset and the ‘New’ dataset. The scores of these tests are recorded and then the model performs the same test *ITER* times. The benchmarking loop also performs data transformation on the raw data in order to test the transformation time as well. After all iterations are over, the average time per loop is recorded in milliseconds. For a good balance of reliable time measuring and computational resource availability, the *ITER* loop was run a total of 10 times. This was repeated for all models and all pipelines.

### A. Naive Bayes

Naive Bayes is a supervised learning technique that is based on the Naive Bayes theorem. Unlike the other models in this report (SVM, Logistic Regression, and the tree based models), Naive Bayes is considered as a generative learning algorithm, focusing to model the distribution of inputs of a given class rather than learning which features are most important for discrimination. Some advantages of Naive Bayes include its simplicity and efficiency, the ability to handle high dimensional data, and its scalability. Some disadvantages include its assumption of conditional independence which causes to to be unrealistic, as well as its subjectability to the “Zero Frequency”, which occurs when a feature does not exist in the training, causing the model to only predict zero.

For the modeling step, ComplementNB(), which is a version of MultinomialNB() that was designed to correct some its severe assumptions and works well with imbalanced data, as well as GaussianNB() was used. Additionally, hyperparameter tuning using GridSearchCV was utilized in order to improve results. ComplementNB() and GaussianNB() have different parameters, so they will be gone over in their respective sections.

1) *Complement Naive Bayes*: Complement Naive Bayes is a variation of Multinomial Naive Bayes that uses the

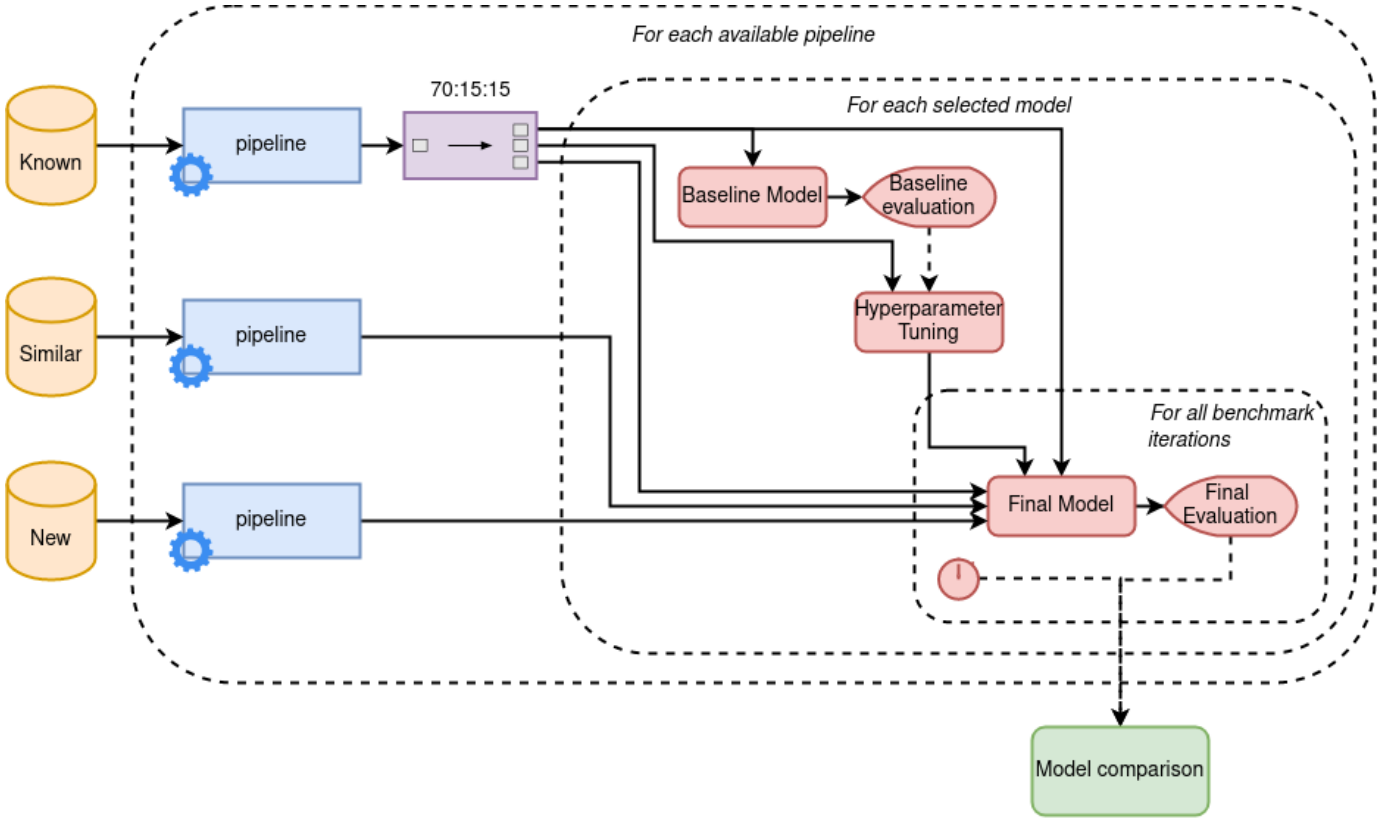


Fig. 3. An overview flow diagram of the project as a whole

complement of each class to calculate weights. As a result, this makes this model more suitable for imbalanced datasets, such as the TRaBID dataset. Complement Naive Bayes has the following parameters: `alpha`, `force_alpha`, `fit_prior`, `class_prior`, and `norm`. `Alpha` is the smoothing parameter, `force_alpha` will set `alpha` to  $1e-10$  if it is too low, `fit_prior` is only used in edge cases with a single classes in the data, `class_prior` is the prior probabilities of the classes but is not used in this implementation, and `norm` performs a second normalization of the weights. The parameter grid is as follows:

```
{
  "alpha": [0.001, 0.01, 0.1, 1, 10, 100, 1000],
  "force_alpha": [True, False],
  "fit_prior": [True, False],
  "norm": [True, False],
}
```

The best performing parameters for each test were the same for each model:

```
{
  "alpha": 0.001,
  "fit_prior": True,
  "force_alpha": True,
  "norm": False
}
```

Featured in Fig. IV are the results from the benchmark function using Complement Naive Bayes model after being going through hyperparameter tuning.

The results seem as expected, where the model performs well on known and similar attacks, while performing lackluster

TABLE IV  
COMPLEMENTNB RESULTS

Pipeline	Dataset	Accuracy	Precision	Recall	F1
scaled	Known	0.940277	0.956870	0.903423	0.929379
	Similar	0.972131	0.969336	0.959866	0.964578
	New	0.794255	0.743730	0.183754	0.294697
corr. scaled	Known	0.944196	0.961817	0.907748	0.934001
	Similar	0.977508	0.975010	0.967912	0.971448
	New	0.796284	0.799373	0.172375	0.283596

for new attacks. The best overall model for this test was on the scaled data, achieving an F1 score of 0.294697. This test, however, had some issues. The original preprocessing used for the benchmark function used a `StandardScaler()` and PCA. However, these methods gave negative values, which cannot be used by this model. As a result, the scaling method was changed to `MinMaxScaler()` and PCA was omitted. But by doing so, the process became different, making comparison a little more difficult, hence why Gaussian Naive Bayes was utilized to create tests that were more inline with the other models.

2) *Gaussian Naive Bayes*: Gaussian Naive Bayes generally performs better on continuous data, especially that which follows the normal distribution. While the data is not necessarily normal, it is continuous due to the `StandardScaler()` that is applied for the benchmark function. Unlike `ComplementNB()`, `GaussianNB()` is able to use negative values, allowing it to

follow the same preprocessing and benchmark tests as the other models. This makes it much easier to make a comparison in terms of overall performance. GaussianNB() has two main parameters: priors, which are the prior probabilities for the classes, and var\_smoothing, which is the portion of the largest variance of all features that is added to variances for stability. In this case, priors was ignored and hyperparameter tuning was only performed based on var\_smoothing. The parameter grid used was  $\{\text{'var\_smoothing': np.logspace(0,-9, num=100)}\}$ , which creates 100 numbers between 1 and  $1e-9$ . Featured in Table V are the results for Gaussian Naive Bayes.

TABLE V  
GAUSSIANNB RESULTS

Pipeline	Dataset	Accuracy	Precision	Recall	F1	Time (ms)
scaled	Known	0.817	0.707	0.991	0.825	9.917
	Similar	0.795	0.661	0.987	0.792	19.195
	New	0.691	0.313	0.268	0.289	50.295
corr, scaled	Known	0.826	0.716	0.995	0.833	4.697
	Similar	0.806	0.672	0.994	0.802	7.989
	New	0.700	0.319	0.251	0.281	28.354
PCA	Known	0.786	0.672	0.992	0.801	15.768
	Similar	0.760	0.625	0.984	0.764	30.440
	New	0.664	0.383	0.711	0.498	83.340
corr, PCA	Known	0.935	0.923	0.927	0.925	12.799
	Similar	0.969	0.943	0.980	0.961	22.165
	New	0.793	0.729	0.186	0.296	41.312

This model still follows the trend of doing well on known and similar attacks while performing mediocre on new attacks. However, this model seems to overfit less and performs significantly better on the new data. Based on what can be compared with ComplementNB(), the overall accuracies and computation time are quite comparable with each other. But, the precision to recall ratio is significantly different. For ComplementNB(), the precision is much greater than the recall, which implies that those models fail to differentiate the attack class well, but is normally correct when it does predict it. On the other hand, for GaussianNB(), the models were more balanced in terms of these metrics. The main difference stemmed from how the models performed on the new attacks.

The best performing test was using all of the features with 95% PCA. This test used a var\_smoothing of  $4.3287e-06$ . In terms of performance with known and similar attacks, it performed around the same, with F1 scores of 0.80096 and 0.76423 respectively. However, the new attacks drastically increased performance in terms of F1 score, achieving an F1 score of 0.49767, which is nearly double of those from the other tests. In terms of precision and recall, the precision of the model increased from around 0.3 to 0.38 while recall had massive increases, going from around 0.25 to 0.7. This means that this model is able to detect attacks better than the other models, but is prone to falsely flagging normal traffic. Despite having a slight accuracy drop, we believe that this is the best performing Naive Bayes model, mainly due to the class imbalance that exists in the data that skews accuracy. The precision of the model is approximately the same as the other tests, but the recall is significantly higher, going from

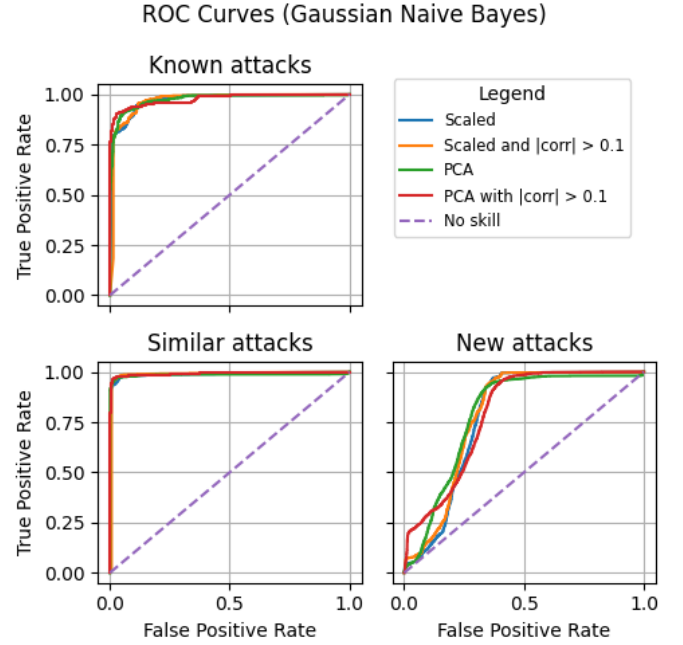


Fig. 4. Gaussian Naive Bayes ROC Curve

around 0.25 to 0.7, implying that it is identifying threats and non threats much better than the others tests. For a problem like this, identifying problems, rather than being cautious, is imperative. Pictured in Fig. 4 are the ROC curves for the known attacks, similar attacks, and new attacks for all of the tests using the GaussianNB() model.

An important thing to note is that the best performing model was using all features with 95% PCA. The removal of features based on correlation was based on the known attacks data. This could imply that some of the features that were removed based on correlation may have been unimportant for known and similar, but were important for new attacks.

### B. Support Vector Machines

Support Vector Machines (SVM) is a supervised machine learning method whose purpose is to separate two classes with an optimal hyperplane that maximizes the margin distances between the support vectors of each class and the hyperplane. This is a discriminative model, which means that it works by comparing the query point with the trained data points and bases the classification off of that. Some of the benefits of working with SVM include better performance on large feature spaces as well as it being memory efficient as it only requires the support vectors when making a classification. Some drawbacks of this model is that it is not easily interpretable and requires a lot of time to train. Theoretically, compared to the other models discussed in this paper, SVM should perform well. However, we will see later that out of all the models tested, SVM performed very poorly. One possible reason for this is that SVM has a tendency to over-fit to the dataset if there are too many samples. The results for SVM can be seen in Table VI.



TABLE VI  
SVM RESULTS

Pipeline	Dataset	Accuracy	Precision	Recall	F1	Time (ms)
scaled	Known	0.998	0.997	0.997	0.997	518.730
	Similar	0.990	0.995	0.981	0.988	1037.263
	New	0.767	0.656	0.009	0.019	2984.307
corr, scaled	Known	0.996	0.997	0.993	0.995	362.891
	Similar	0.989	0.998	0.975	0.986	749.064
	New	0.766	0.383	0.003	0.006	2147.515
PCA	Known	0.997	0.996	0.997	0.996	689.647
	Similar	0.991	0.997	0.980	0.989	1514.054
	New	0.767	0.620	0.008	0.016	3798.192
corr, PCA	Known	0.996	0.999	0.993	0.996	428.190
	Similar	0.989	0.999	0.972	0.985	830.623
	New	0.766	0.586	0.004	0.009	2294.048

We used SciKit-Learn’s SVC classifier, which is based on the ‘LibSVM’ implementation [8], [9]. Using this requires us to select some parameters for modelling. The main parameters we used for training are as follows.

The `kernel` parameter controls which kernel the SVC classifier should use. The kernels supported by ‘LibSVM’ are the linear kernel, polynomial kernel, RBF kernel, and sigmoid kernel. The benefit of using a specific kernel is how well the shape of the data fits to the given function. If there is a clear linear separation between the two classes, the best kernel will be the linear kernel, and so on.

There is a concept in SVM known as softmax margin. A softmax margin means that there can be some error in the margin’s decision boundary, allowing for more flexible models. The `C` parameter controls the strength of this regularization.

Finally, we have the `gamma` parameter. This parameter takes on different meanings depending on the kernel being used. If the RBF kernel is used, then this parameter controls the radius of influence of training samples. If the polynomial kernel is used, this parameter controls the domination of higher degree terms. If the sigmoid kernel is used, this parameter controls the steepness of the sigmoid curve. This parameter is not used in the linear kernel. [8]

The parameter grid is therefore designed to test different combinations of each of these parameters in order to find the optimal model for the problem. The specific values for our parameter grid are shown as the following

```
{
  "C": [0.01, 0.1, 1, 10, 100, 1000],
  "kernel": ["poly", "rbf", "sigmoid",
             "linear"],
  "gamma": ["scale", "auto"],
}
```

Note, the ‘scale’ value for `gamma` means that the value will be set to  $1/(n \times \sigma_X^2)$ , while setting it to ‘auto’ will set the value to simply  $1/n$ . Where  $n$  is the number of features and  $\sigma_X^2$  is the variance of the input features. The optimal parameters for SVM based on given parameter grid can be seen in Table VII.

Given the results in Table VI, we can see that even the best performing model only has a F1 score of 0.019 on the ‘New’ dataset. One thing to note is that the best performing pipeline

TABLE VII  
SVM BEST HYPERPARAMETERS

Pipeline	scaled	corr, scaled	PCA	corr, PCA
C	100	1000	100	1000
gamma	scale	scale	auto	auto
kernel	rbf	rbf	rbf	rbf

ROC Curves (SVM)

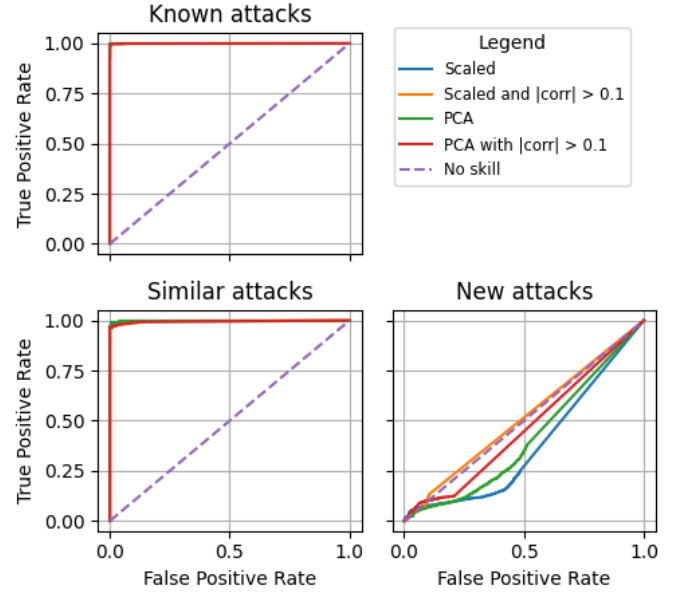


Fig. 5. SVM ROC curves for all pipelines and each of the datasets tested. All models performed poorly except for the ‘corr, scaled’ model, which just barely exceeds the no skill line

in SVM was the pipeline that had the most number of features, i.e. the ‘scaled’ pipeline, which used all original features passed through a standard scaler. In terms of detection time, SVM performed the worst of all models and pipelines tested in this project. The ROC curves for this model’s performance on all datasets and pipelines can be seen in Fig. 5.

Taking a close look at the ROC curves and the result table is that even the ‘testing’ partition of the ‘Known’ dataset has a nearly perfect prediction score across all pipelines, with the lowest F1 score being 0.995 and the lowest accuracy being 0.996. Similarly, with the ‘Similar’ dataset, we can see that it too has very high scores. This might be indicative of the model overfitting to the ‘Known’ dataset, and not being able to generalize to the concept of a harmful ‘attack’ network event.

### C. Logistic Regression

Logistic regression is a statistical model used for classification problem. It is mainly used for predicting binary outcomes where the dependent variable is categorical. It is very easy to implement and interpret compared to other models. It saves computation costs and is efficient in scaling with large datasets. Moreover, it gives us probabilistic estimates which aid in efficient decision making but, it also has its drawbacks.

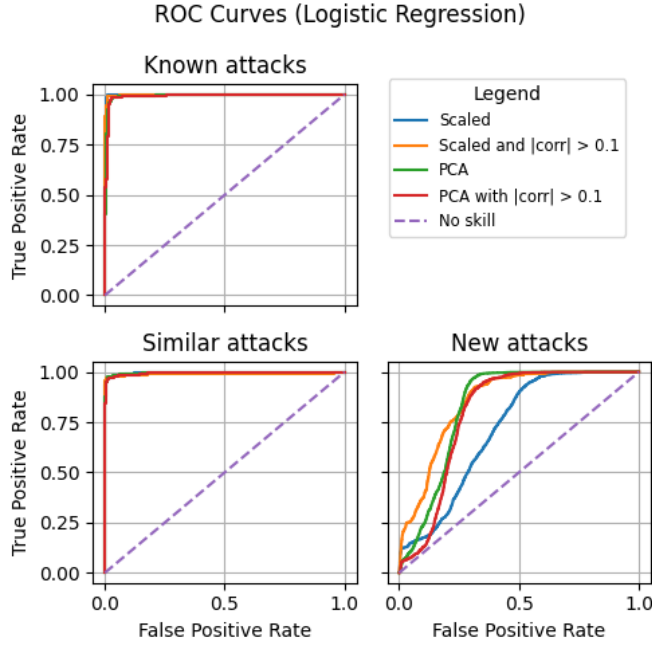


Fig. 6. Logistic Regression ROC Curve

It usually assumes a linear relationship between the features and target features. It is also sensitive to outliers and can affect model performance. The results for Logistic Regression can be seen in Table VIII.

The “Known attacks” ROC curves are almost perfect in this case; TPR is almost 1.0 FPR is kept low. The preprocessing steps or approaches, like “Scaled”, “Scaled and  $|corr| > 0.1$ ”, “PCA”, “PCA with  $|corr| > 0.1$ ”, and “No skill” are all listed on the left side of the legend. Just like for “Known attacks”, “Similar attacks” ROC curves are very accurate as well. The TPR for all techniques is extremely high, implying that the method can reliably detect attacks which are similar to the ones that are already found. For the new attacks we can see that there is more variation with performance which is lower than known and similar attacks. “Scaled” and “Scaled and  $|corr| > 0.1$ ” methods perform better and show a high TPR at lower FPR when compared to “PCA” and “PCA with  $|corr| > 0.1$ ”.

TABLE VIII  
LOGISTIC REGRESSION RESULTS

Pipeline	Dataset	Accuracy	Precision	Recall	F1	Time (ms)
scaled	Known	0.993	0.990	0.995	0.992	6.337
	Similar	0.980	0.995	0.954	0.974	18.706
	New	0.774	0.845	0.044	0.084	36.033
corr, scaled	Known	0.990	0.983	0.993	0.988	11.030
	Similar	0.984	0.988	0.972	0.980	15.148
	New	0.785	0.822	0.101	0.181	24.301
PCA	Known	0.980	0.974	0.979	0.977	16.595
	Similar	0.983	0.983	0.972	0.978	22.989
	New	0.767	0.533	0.042	0.078	72.396
corr, PCA	Known	0.978	0.974	0.977	0.975	7.628
	Similar	0.982	0.981	0.973	0.977	13.880
	New	0.767	0.546	0.033	0.061	30.205

The scaled pipeline shows us that there is a high performance on known and similar datasets with accuracy’s 0.993 and 0.990. however we can see a performance drop on the new dataset with an accuracy of 0.774. This shows us that the model is struggling with generalization in spite of strong performance on familiar data. The corr, scaled pipeline shows us that it is performing very well on similar and known attacks with accuracy’s of 0.994 and 0.985. We can see a slight improvement for the new attacks dataset with the accuracy of 0.788 which suggests that correlation feature selection might enhance accuracy when compared to single scaling. The PCA pipeline is performing well on known and similar attacks datasets with 0.994 and 0.987 as the accuracy’s. The performance on the new attacks has dropped to 0.767 which suggests that PCA struggles to identify the new attacks. The corr, PCA pipeline shows high performance on Known (0.993) and Similar (0.978) datasets. However, it performs the worst on the new attacks with an accuracy of 0.767. The scaled and corr, scaled pipelines are running at a faster rate with time ranging from 6 to 38 ms. The PCA and corr, PCA pipelines are taking a longer time to run which is between 15 to 34 ms. This tells us that PCA might need complex computations and does not help significantly with the model on new data. We can see that the precision and recall drops significantly for new attacks which indicates that it is struggling to identify true positives. It also consists of lowest precision of 0.546 which lowest when compared to known and similar attacks.

#### D. Random Forest

Random Forest is an ensemble learning technique that is particularly useful in threat detection projects because of its scalability, robustness, versatility, and feature importance insights. It is utilized for both classification and regression problems. Using randomized subsets of the training set, it builds numerous decision trees and combines their results to improve prediction accuracy and manage overfitting. Because of this, Random Forest works well with threat detection datasets that are complicated and variable. Hyper parameters like `n_estimators` (number of trees), `max_depth` (maximum depth of each tree), `min_samples_leaf` (minimum samples required to be at a leaf node), and `split_samples` (minimum samples required to split an internal node) can be used to fine tune the model’s performance. By balancing the biasvariance tradeoff, these parameters increase the model’s efficacy and accuracy in detecting threats.

Table IX shows the outcomes of using a Random Forest classifier to analyze several machine learning pipelines and quickly detect network threats on the TRaBID dataset. The performance measurements for each pipeline include accuracy, precision, recall, F1 score, and calculation time (measured in milliseconds). The dataset is divided into three categories: known, similar, and new threats. With a computation time of 140.314 ms, the Scaled pipeline achieves flawless scores across all measures, demonstrating excellent performance against known threats. With a computation time of 242.981 ms, it retains high accuracy (0.987), precision (1.000), recall (0.968), and F1 score (0.984) for similar threats. With an

accuracy of 0.767, precision of 1.000, recall of 0.003, F1 score of 0.006, and the maximum computation time of 622.699 ms. New threats, on the other hand, show a marked decline in performance.

TABLE IX  
RANDOM FOREST RESULTS

Pipeline	Dataset	Accuracy	Precision	Recall	F1	Time (ms)
scaled	Known	1.000	1.000	1.000	1.000	140.314
	Similar	0.987	1.000	0.968	0.984	242.981
	New	0.767	1.000	0.003	0.006	622.699
corr, scaled	Known	1.000	1.000	1.000	1.000	68.079
	Similar	0.989	1.000	0.971	0.985	111.215
	New	0.767	0.923	0.005	0.011	308.599
PCA	Known	0.994	0.992	0.995	0.993	157.932
	Similar	0.989	0.995	0.977	0.986	232.806
	New	0.767	0.578	0.012	0.023	525.660
corr, PCA	Known	0.997	0.996	0.997	0.996	84.760
	Similar	0.990	0.997	0.978	0.988	135.499
	New	0.767	0.629	0.006	0.012	272.315

With the shortest computation time of 68.079 ms, the “Correlation and Scaled” pipeline obtains flawless scores for known threats. With an accuracy of 0.989, precision of 1.000, recall of 0.971, F1 score of 0.985, and computation time of 111.215 ms, it outperforms the Scaled pipeline for Similar threats. It performs better than the Scaled pipeline when it comes to New threats, with a computation time of 308.599 ms, accuracy of 0.767, precision of 0.923, recall of 0.005, and F1 score of 0.011. With an accuracy of 0.994, precision of 0.992, recall of 0.995, F1 score of 0.993, and computation time of 157.932 ms, the PCA pipeline performs marginally worse for known threats. Performance drops much more for similar threats: 232.806 ms of computation time, 0.989 accuracy, 0.995 precision, 0.978 recall, and 0.986 F1 score. With an accuracy of 0.767, precision of 0.578, recall of 0.012, F1 score of 0.023, and computation time of 525.660 ms, the performance of the New threats is comparable to that of the Scaled pipeline.

With an accuracy of 0.997, precision of 0.996, recall of 0.997, F1 score of 0.996, and computation time of 84.760 ms, the Correlation and PCA pipeline performs well for known threats. With an accuracy of 0.990, precision of 0.997, recall of 0.978, F1 score of 0.988, and computation time of 135.499 ms, it performs well against similar threats. For New threats, performance is poor, with an accuracy of 0.767, precision of 0.629, recall of 0.006, F1 score of 0.012, and a computation time of 272.315ms.

The ROC curves can be seen in Fig. 7. ROC (Receiver Operating Characteristic) curves that indicate how well a Random Forest classifier performs in threat detection for three different scenarios: known assaults, similar attacks, and new attacks. Before the classifier is trained, various feature scaling and selection techniques are represented by each curve. The percentage of genuine positives that the model correctly detected is known as the True Positive Rate (TPR), sometimes referred to as sensitivity or recall. On the other hand, the percentage of actual negatives that the model mistakenly

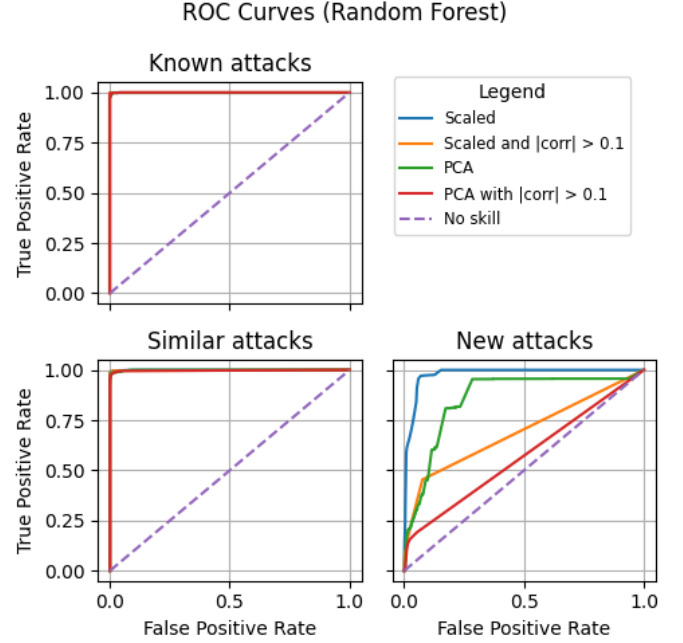


Fig. 7. Random Forest ROC Curve

labeled as positives is known as the False Positive Rate (FPR). Plotting TPR against FPR at different threshold values, the ROC curve shows that higher performance is indicated by a closer distance to the top-left corner. “PCA” applies Principal Component Analysis to reduce dimensionality as discussed in section IV; “PCA with  $|corr| > 0.1$ ” combines PCA with the correlation threshold; “Scaled” indicates standard scaling; “Scaled and  $|corr| > 0.1$ ” retains only features with an absolute correlation greater than 0.1; and “No skill” represents a baseline model making random predictions. All techniques yield almost flawless identification for known attacks, demonstrating the effectiveness of the classifier when it has encountered similar patterns during training. Excellent performance is also demonstrated by similar attacks, indicating good generalization to variations of well-known methods. Performance varies for new attacks; the “Scaled” technique works best, suggesting that resilience against new threats depends on maintaining all characteristics with appropriate scaling. While approaches involving correlation thresholds perform less well and may exclude significant features, “PCA” performs quite well. These results emphasize the significance of thorough feature representation and suitable preprocessing, with scaling without further filtering offering the best generalization to novel threats and therefore improving threat detection systems’ efficacy.

#### E. XGBoost

XGBoost is a popular supervised tree based algorithm that is popularly used in classification problems. This model is an example of an ensemble technique which follows the boosting principle. XGBoost is beneficial for several reasons. It tends to deliver a relatively high accuracy. It is also known for being very efficient when making predictions making them



suitable for real time applications. XGBoost is not heavily impacted by outliers and is capable of handling both large and imbalanced datasets. The model is resistant to underfitting by default and has built-in parameters that can help to prevent overfitting. Along with the benefits, the model also has some drawbacks. One case is where despite having parameters that help it prevent overfitting, because of the model's nature, it can still be susceptible to it especially if not properly tuned. XGBoost's performance can also be a hit or miss on small or sparse datasets. This model can also take a while to train on large datasets since it is memory intensive. It can take up a good chunk of space within the RAM. Since this model has a lot of parameters in it, tuning the model can take time.

In this case, we used the XGBClassifier available in the xgboost library to train this model. Generally speaking, this model should perform really well. In many other cases, it is generally considered the best performing model out of all. Here, the XGBoost model mostly performed as was expected for each test case. While this model was not the worst performing one, we can tell that it certainly did not perform the best. The results for XGBoost can be seen in Table X.

TABLE X  
XGBCLASSIFIER RESULTS

Pipeline	Dataset	Accuracy	Precision	Recall	F1	Time (ms)
scaled	Known	1.000	1.000	1.000	1.000	39.135
	Similar	0.989	0.999	0.973	0.986	51.636
	New	0.762	0.276	0.010	0.019	126.805
corr, scaled	Known	1.000	1.000	1.000	1.000	22.749
	Similar	0.991	0.999	0.978	0.988	46.507
	New	0.780	0.751	0.089	0.159	122.587
PCA	Known	0.996	0.994	0.996	0.995	89.676
	Similar	0.991	0.996	0.982	0.989	105.654
	New	0.767	0.602	0.017	0.034	169.958
corr, PCA	Known	0.997	0.996	0.996	0.996	76.115
	Similar	0.990	0.995	0.980	0.987	88.797
	New	0.767	0.692	0.011	0.021	138.890

Based on the result, we can see that the model performed the best on the 'new' dataset through the correlation scaled pipeline. The best f1-score obtained was around 0.159. This result is better than some of the models, but still not as good in performance as others. This instance of the model on all versions training the 'new' data also had the fastest training time at around 122 milliseconds.

To obtain the best possible results for every test case, we utilized hyperparameter tuning using GridSearchCV to train the model using the best possible parameter combination. The following parameter grid was utilized in to perform the hyperparameter tuning:

```
{
  "max_depth": [10, 20, 30, 40, 50],
  "n_estimators": [100, 200, 400, 800],
  "learning_rate": [0.01, 0.1, 0.2],
  "gamma": [0.1, 0.5, 1],
  "scale_pos_weight": [0.1, 1, 5, 10]
}
```

The max\_depth parameter identifies the maximum depth of the tree. It controls the maximum number of nodes the exist

in the tree from the root to the leaf node. This parameter is responsible for the complexity of the model. If the max\_depth value is set to being too high, there is the risk of overfitting. The higher the depth, the longer it will take for the model to train.

The n\_estimators focuses on the number of trees built or added to the model behind the scenes. Every subsequent tree learns from the previous trees mistakes. This higher the number of estimators, the more complex the model becomes. If the number of trees is too high, it can cause overfitting. The speed at which the model can train and predict the results also gets slower if the n\_estimators is too high.

The learning\_rate is responsible for scaling the weights to determine the step size to take in each boosting round. It determines how quickly the model converges to its optimal point. Typically a lower learning rate will require more iterations for the model to reach convergence. The size of the learning\_rate can impact the accuracy of the model. Typically if there are a high number of n\_estimators, then a lower learning rate yields to a model generalizing better and yielding to more accurate results while also reducing the risk of overfitting.

The gamma parameter represents the minimum loss reduction needed for a split to occur at a leaf node. Typically the higher the gamma value, the more the tree pruning there is. This prevents any unnecessary splits from, occurring. Increasing gamma also reduces the model complexity which can help reduce the risks of overfitting. Gamma acts as the regularization parameter in XGBoost which works to prevent overfitting from occurring.

The scale\_pos\_weight addresses the balance between positive and negative classes. In other words, it handles class imbalance issues in a dataset with binary classification. Using this parameter can help to increase the accuracy or f1-score and try to prevent super skewed results from being obtained if the data is imbalanced.

Figure 8 provides the ROC curves showing the distribution of classification for all datasets.

Based on the curves for 'new attacks', we can tell the model performed as expected. All the cases involving PCA performed worse than any of the scaled versions. The red curve representing the PCA with correlation pipeline shows more variance in results compared with the other cases. For every other dataset in each pipeline, the f1-score is never below 0.9. For every instance of the 'new' dataset, the result, again, appears to be the best in the pipeline involving scaled with correlation and also has the fastest computation time. For every other pipeline, the f1-score is below 0.05. XGBoost's overall performance was generally reasonable, but not as good as how Naive Bayes or Logistic Regression performed.

## VI. EVALUATING RESULTS AND COMPARING

Now that we have all the various models trained, we can run all models on the same system to perform a comparative benchmark. Since false positives and false negatives are both bad for our application, we will primarily compare the models using F1 score. At the same time, we also need to find the best model in terms of detection time. We will compare the

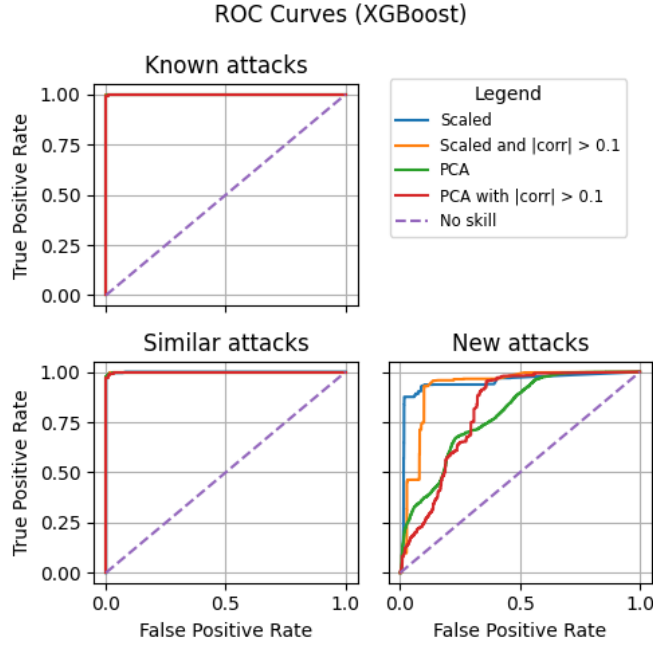


Fig. 8. ROC curves for XGBoost

models in these regards using the ‘New’ dataset as this is most important dataset for making sure that the selected model performs well in the real world.

The tables comparing best models in terms of F1 score can be seen in Table XI. The best performing model in terms of F1 score was Gaussian Naive Bayes using all features with 95% PCA.

TABLE XI  
TOP 5 MODELS BY F1 ON NEW DATA

Rank	Model	Pipeline	F1
1	GaussianNB	PCA	0.498
2	GaussianNB	corr, PCA	0.296
3	GaussianNB	scaled	0.289
4	GaussianNB	corr, scaled	0.281
5	Logistic Regression	corr, scaled	0.181

The tables comparing best models in terms of detection time can be seen in Table XII. The best performing model in terms of speed was Logistic Regression using scaled features that had a correlation of greater than 0.1.

TABLE XII  
TOP 5 MODELS BY TIME (MS) ON NEW DATA

Rank	Model	Pipeline	Time (ms)
1	Logistic Regression	corr, scaled	24.301
2	GaussianNB	corr, scaled	28.354
3	Logistic Regression	corr, PCA	30.205
4	Logistic Regression	scaled	36.033
5	GaussianNB	corr, PCA	41.312

By looking at both of these tables, it is clear that both

Logistic Regression and Gaussian Naive Bayes perform the best of all the models in this project. Gaussian Naive Bayes far exceeds all models in terms of F1 score, even when using any of the available pipelines. Logistic regression and Gaussian Naive Bayes are both equally represented in the top 5 in terms of detection time with all models approximately around 30ms.

## VII. DEPLOYMENT

We have two potential deployment strategies for this project. First, we could make this project available online for research purposes. Second, we could select the ‘best’ model from this project and make it available as an end user firewall application.

For the first possible deployment solution, we can host all model pickle files to platforms such as HuggingFace. The use case for this deployment solution is that it will allow other researchers and developers to use our models when performing studies. These online platforms will allow easy scaling to a vast number of users across the world as long as the compute is available. One drawback to this approach is that there would be a continuous cost to the account owners.

For the second approach, we can create a user application in Python that is installed to the user’s network. This application will come bundled with the best model from this study, Gaussian Naive Bayes, and will apply the best pipeline to live network traffic. This live network traffic will be evaluated by the model and will alert the user in case of detection of malicious network activity. Scaling this will be relatively simple as each user will need to download and install the application, only providing their own computation resources.

For either of these deployment strategies, we would first need to improve the models and provide the model saved to disk. This can be done with the `pickle` package in Python.

## VIII. CONCLUSION

Based on our poor results, we believe that the majority of models are massively overfitting. For almost every model, every statistic for the known and similar datasets are 1, or very close to it, while the F1 scores for new attacks are extremely low. This is likely due to the fact that the new attacks data is very dissimilar with the known and similar attacks. As a result, the models are unable to properly predict new attacks. The only model that seemed to combat this was Naive Bayes. This could be because Naive Bayes is considered a generative model whereas the other models are considered discriminative. Discriminative models identifies decision boundaries and categorizes based on that. On the other hand, generative models learns the probability distribution of the data.

This project had a few limitations and things that could be improved on. The original known attacks data had over 28 million rows while we only used 100,000. This was due to a lack of sufficient hardware and computation power. Additionally, given that every model was unable to perform well, we were unable to significantly link the known and similar attacks to the new attacks.

## REFERENCES

- [1] “DDoS threat report 2023 Q4,” 2023. [Online]. Available: <https://blog.cloudflare.com/ddos-threat-report-2023-q4/>
- [2] “2023 data breach investigations report,” 2023. [Online]. Available: <https://www.verizon.com/business/resources/Tf14/reports/2023-data-breach-investigations-report-dbir.pdf>
- [3] S. Bagui, E. Kalaimannan, S. Bagui, D. Nandi, and A. Pinto, “Using machine learning techniques to identify rare cyber-attacks on the unsw-nb15 dataset,” *Security and Privacy*, vol. 2, no. 6, p. e91, 2019.
- [4] N. S. Bhati and M. Khari, “An ensemble model for network intrusion detection using adaboost, random forest and logistic regression,” in *Applications of Artificial Intelligence and Machine Learning: Select Proceedings of ICAAAIML 2021*. Springer, 2022, pp. 777–789.
- [5] H.-G. Zhou *et al.*, “Using immune algorithm to optimize anomaly detection based on svm,” in *2006 International Conference on machine learning and cybernetics*. IEEE, 2006, pp. 4257–4261.
- [6] M. Tavallae, E. Bagheri, W. Lu, and A. A. Ghorbani, “A detailed analysis of the kdd cup 99 data set,” in *2009 IEEE Symposium on Computational Intelligence for Security and Defense Applications*, 2009, pp. 1–6.
- [7] E. K. Viegas, A. O. Santin, and L. S. Oliveira, “Toward a reliable anomaly-based intrusion detection in real-world environments,” *Computer Networks*, vol. 127, pp. 200–216, 2017.
- [8] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [9] C.-C. Chang and C.-J. Lin, “Libsvm: a library for support vector machines,” *ACM transactions on intelligent systems and technology (TIST)*, vol. 2, no. 3, pp. 1–27, 2011.